

1_exploratory_data_analysis

November 2, 2020

1 Exploratory Data Analysis

```
[1]: # Install watermark to print system and hardware information
# %conda install watermark
%load_ext watermark
%watermark -v -m -p numpy,pandas,matplotlib,seaborn
```

```
CPython 3.8.6
IPython 7.18.1
```

```
numpy 1.19.2
pandas 1.1.3
matplotlib 3.3.2
seaborn 0.11.0
```

```
compiler      : Clang 10.0.1
system        : Darwin
release       : 19.6.0
machine       : x86_64
processor      : i386
CPU cores     : 16
interpreter    : 64bit
```

1.1 Import Data

The first step of data preparation is to import data. We use `pandas`'s `read_csv()` to import data and take care of data types, true/false values and missing values.

```
[1]: import numpy as np
import pandas as pd
```

```
[2]: def import_dataset(filename):
    bank_mkt = pd.read_csv(filename,
                           na_values=["unknown", "nonexistent"],
                           true_values=["yes", "success"],
                           false_values=["no", "failure"])

    # Treat pdays = 999 as missing values
    bank_mkt["pdays"] = bank_mkt["pdays"].replace(999, pd.NA)
```

```

# Convert types, "Int64" is nullable integer data type in pandas
bank_mkt = bank_mkt.astype(dtype={"age": "Int64",
                                   "job": "category",
                                   "marital": "category",
                                   "education": "category",
                                   "default": "boolean",
                                   "housing": "boolean",
                                   "loan": "boolean",
                                   "contact": "category",
                                   "month": "category",
                                   "day_of_week": "category",
                                   "duration": "Int64",
                                   "campaign": "Int64",
                                   "pdays": "Int64",
                                   "previous": "Int64",
                                   "poutcome": "boolean",
                                   "y": "boolean"})

# reorder categorical data
bank_mkt["education"] = bank_mkt["education"].cat.
→reorder_categories(["illiterate", "basic.4y", "basic.6y", "basic.9y", "high.
→school", "professional.course", "university.degree"], ordered=True)
bank_mkt["month"] = bank_mkt["month"].cat.reorder_categories(["mar", "apr", "
→jun", "jul", "may", "aug", "sep", "oct", "nov", "dec"], ordered=True)
bank_mkt["day_of_week"] = bank_mkt["day_of_week"].cat.
→reorder_categories(["mon", "tue", "wed", "thu", "fri"], ordered=True)
return bank_mkt

```

```
[3]: bank_mkt = import_dataset("../data/BankMarketing.csv")
```

1.2 Exploratory Data Analysis

Exploratory Data Analysis is a process to explore the dataset with no assumptions or hypothesis. The objective is to give us enough insights for the future work.

There are many visualization libraries in Python. Pandas has its own plot API based on matplotlib and we will also use Seaborn and Altair. Seaborn is a Python data visualization library based on matplotlib. It provides a high-level interface for drawing attractive and informative statistical graphics. Altair is a declarative statistical visualization library for Python, based on Vega and Vega-Lite. Both libraries provide easy to use APIs and produce beautiful graphs.

```

[4]: import altair as alt
import matplotlib.pyplot as plt
# cosmetic options for matplotlib
plt.style.use("seaborn")
plt.rcParams["figure.figsize"] = (6.4, 4.8)
plt.rcParams["figure.dpi"] = 300
plt.rcParams["axes.titleweight"] = "bold"

```

```
plt.rcParams["axes.titlepad"] = 10.0
plt.rcParams["axes.titlelocation"] = "left"
from IPython.display import set_matplotlib_formats
set_matplotlib_formats("svg")
import seaborn as sns
```

Let's first inspect the outcome distribution. As we can see below, the dataset is imbalanced. With 41188 rows of data, only 11.2% have positive outcome.

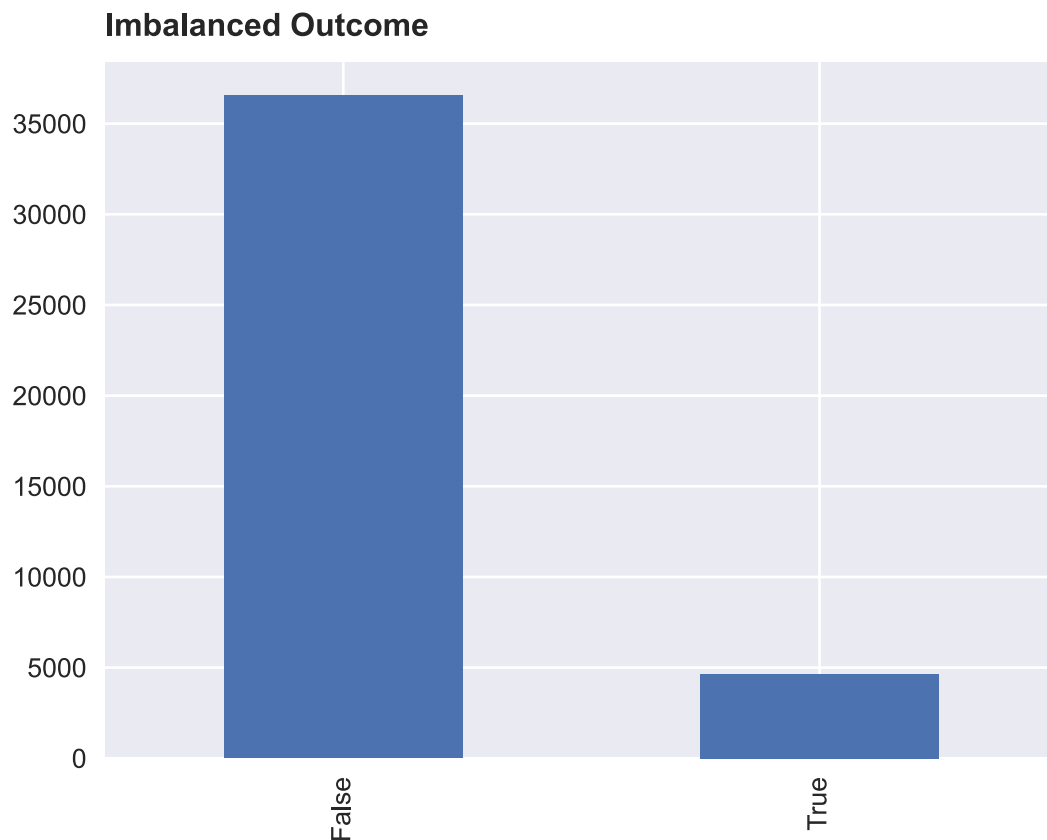
```
[5]: bank_mkt["y"].count()
```

```
[5]: 41188
```

```
[6]: bank_mkt["y"].sum()/bank_mkt["y"].count()
```

```
[6]: 0.11265417111780131
```

```
[7]: y_count = bank_mkt["y"].value_counts().plot(kind = "bar", title="Imbalanced_↵
↵Outcome")
```



Using `info()` we can see that most of features concerning the client are categorical/boolean type. And some fields such as `job`, `marital`, `education`, etc. are missing.


```
[8]: bank_mkt.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   41188 non-null  Int64
1   job                   40858 non-null  category
2   marital               41108 non-null  category
3   education             39457 non-null  category
4   default               32591 non-null  boolean
5   housing               40198 non-null  boolean
6   loan                  40198 non-null  boolean
7   contact               41188 non-null  category
8   month                 41188 non-null  category
9   day_of_week           41188 non-null  category
10  duration              41188 non-null  Int64
11  campaign              41188 non-null  Int64
12  pdays                1515 non-null   Int64
13  previous              41188 non-null  Int64
14  poutcome              5625 non-null   boolean
15  emp.var.rate          41188 non-null  float64
16  cons.price.idx         41188 non-null  float64
17  cons.conf.idx          41188 non-null  float64
18  euribor3m             41188 non-null  float64
19  nr.employed            41188 non-null  float64
20  y                     41188 non-null  boolean
dtypes: Int64(5), boolean(5), category(6), float64(5)
memory usage: 4.0 MB
```

1.2.1 Missing values

By checking the number of missing values, we can see nearly all client do not have `pdays` and `poutcome`. 20% of the clients do not have information of `default`.

```
[9]: na = bank_mkt.isna().sum()
na_nonzero = na[na != 0]
na_perc = na_nonzero/bank_mkt.y.count()
na_bar = na_perc.plot.bar(title="Percentage of Missing Values")
```




1.2.2 Client Data

Let's start with basic client data.

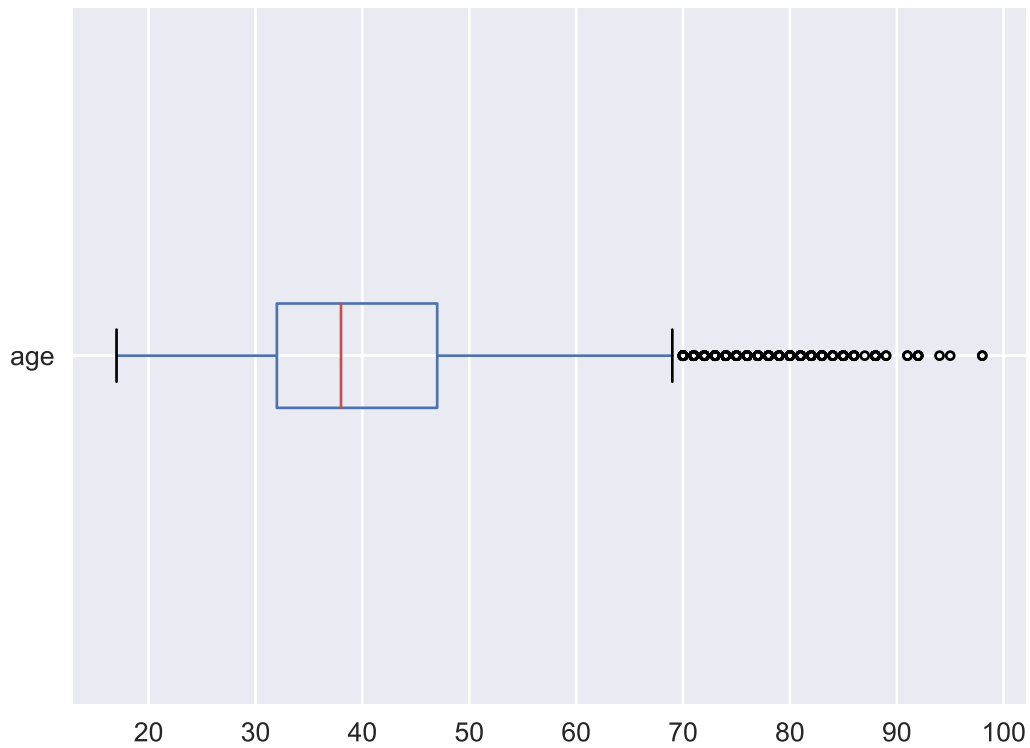
Most of the clients's age are between 32 to 47 while there are some outlier cases beyond 70. This may imply that we should choose standardization for scaling since it's more tolerant for outliers.

```
[10]: age_hist = bank_mkt["age"].plot.hist(title="Age Histogram")
```




```
[11]: age_box = bank_mkt["age"].plot.box(vert=False, sym=".", title="Age_↪Distribution")
```

Age Distribution

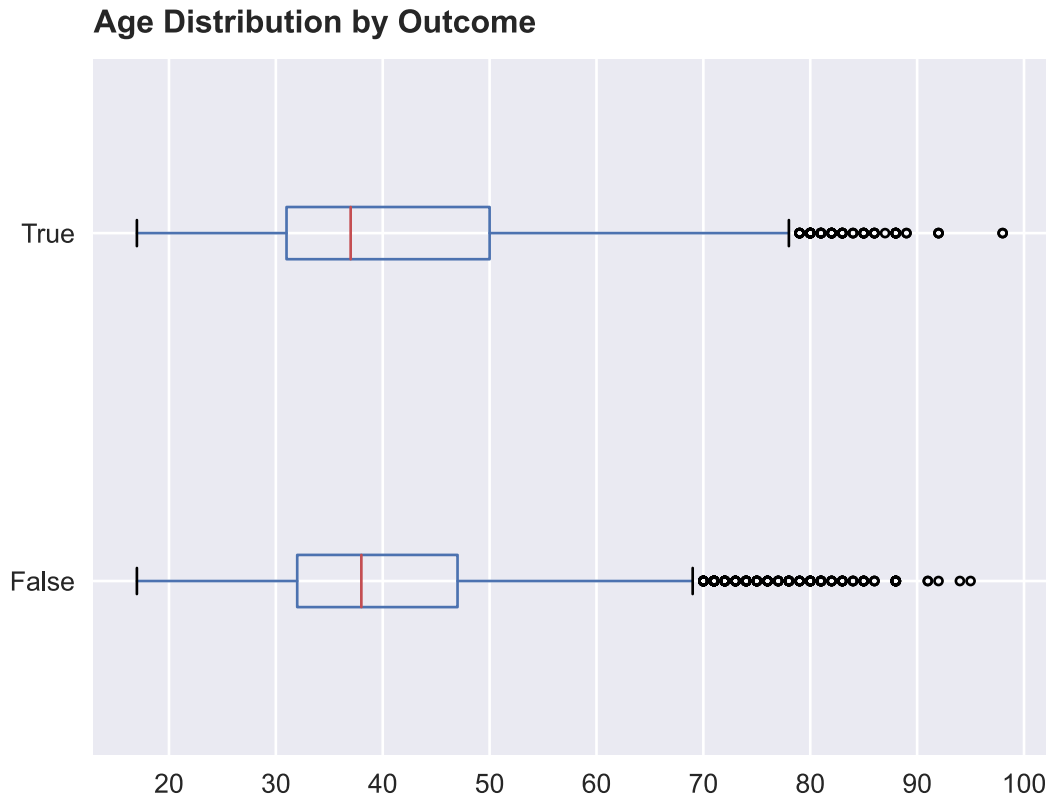


From the graph below we can see that the age distribution in the true outcome group has lower median age but is more skewed toward an slightly older population.

```
[12]: age_y = bank_mkt[["age", "y"]].pivot(columns="y", values="age")
      age_hist_outcome = age_y.plot.hist(alpha=0.9, legend=True, title="Age Histogram_
      ↳by Outcome")
```




```
[13]: age_box_outcome = age_y.plot.box(vert=False, sym=".", title="Age Distribution_
    ↳by Outcome")
```

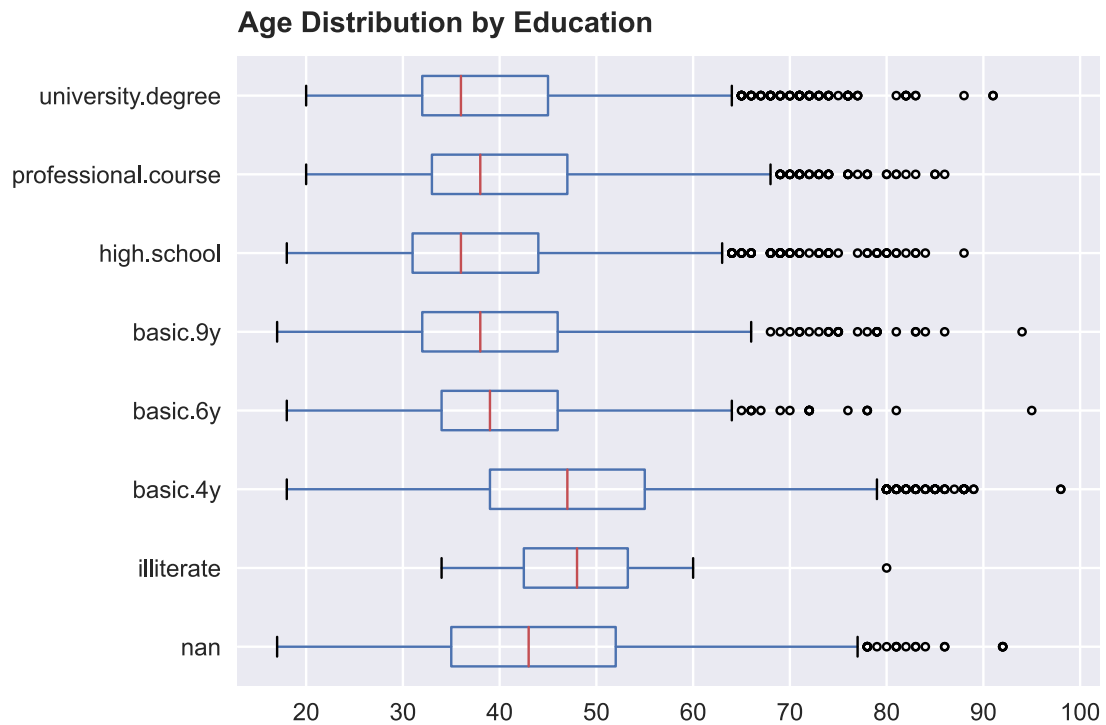



We can also inspect the relationship between age and other categorical values.

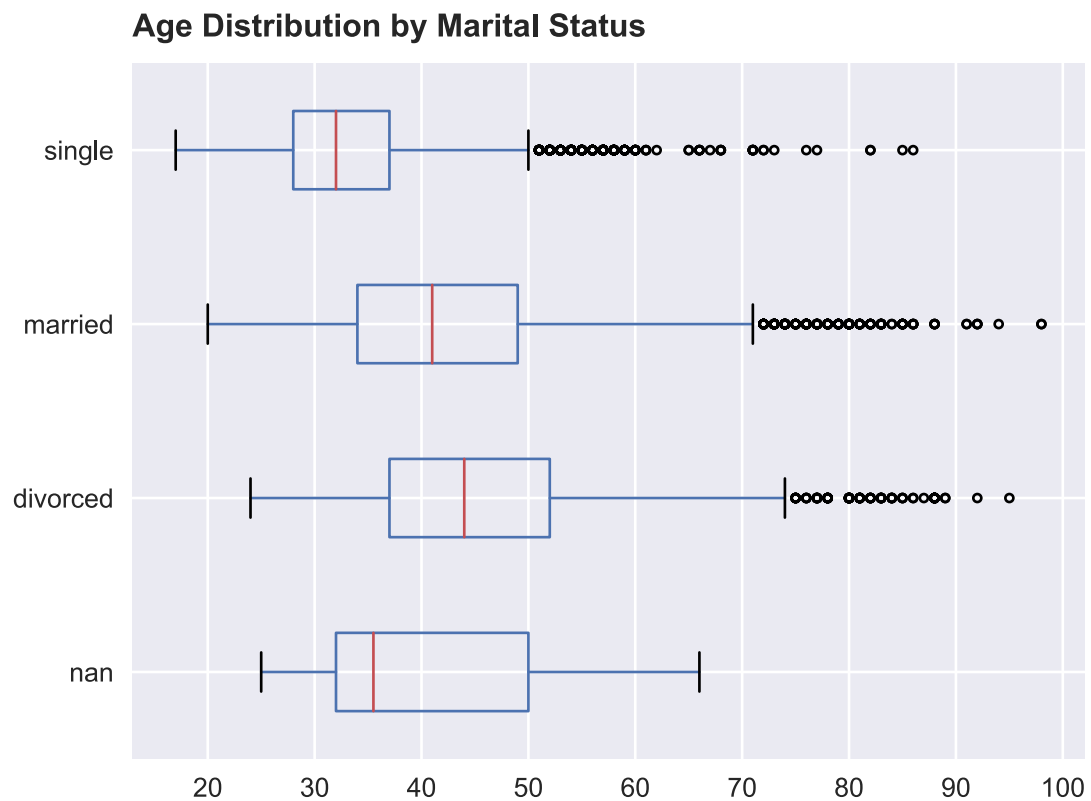
```
[14]: age_job = bank_mkt[["age", "job"]].pivot(columns="job", values="age")
      age_job_box = age_job.plot.box(vert=False, sym=".", title="Age Distribution by_
      ↪Job")
```



```
[15]: age_education = bank_mkt[["age", "education"]].pivot(columns="education",
↪values="age")
age_education_box = age_education.plot.box(ver=False, sym=".", title="Age
↪Distribution by Education")
```




```
[16]: age_marital = bank_mkt[["age", "marital"]].pivot(columns="marital",
    ↪values="age")
age_marital_box = age_marital.plot.box(verte=False, sym=".", title="Age
    ↪Distribution by Marital Status")
```




```
[17]: age_default = bank_mkt[["age", "default"]].pivot(columns="default",
    ↪values="age")
age_default_box = age_default.plot.box(vert=False, sym=".", title="Age
    ↪Distribution by Default")
```

Age Distribution by Default




```
[18]: age_housing = bank_mkt[["age", "housing"]].pivot(columns="housing",
    ↪ values="age")
age_housing_box = age_housing.plot.box(vert=False, sym=".", title="Age
    ↪ Distribution by Housing")
```

Age Distribution by Housing



```
[19]: age_loan = bank_mkt[["age", "loan"]].pivot(columns="loan", values="age")
age_loan_box = age_loan.plot.box(vert=False, sym=".", title="Age Distribution_
↳by Loan")
```

Age Distribution by Loan



We can then turn to job, education and other categorical data to see their relationship to the outcome.

```
[20]: def explore_cat(df, feature):
    df = df.copy()
    if pd.api.types.is_categorical_dtype(df[feature]):
        df[feature] = df[feature].cat.add_categories('unknown')
        df[feature] = df[feature].fillna("unknown")
    feature_true = df[[feature, "y"]].groupby([feature]).sum().y.rename("True")
    feature_total = df[[feature, "y"]].groupby([feature]).count().y.
    →rename("Total")
    feature_false = feature_total - feature_true
    feature_false = feature_false.rename("False")
    feature_true_rate = feature_true / feature_total
    feature_true_rate = feature_true_rate.rename("True Percentage")
    explore_df = pd.concat([feature_true, feature_false, feature_total,
    →feature_true_rate], axis=1).reset_index()
    return explore_df
```

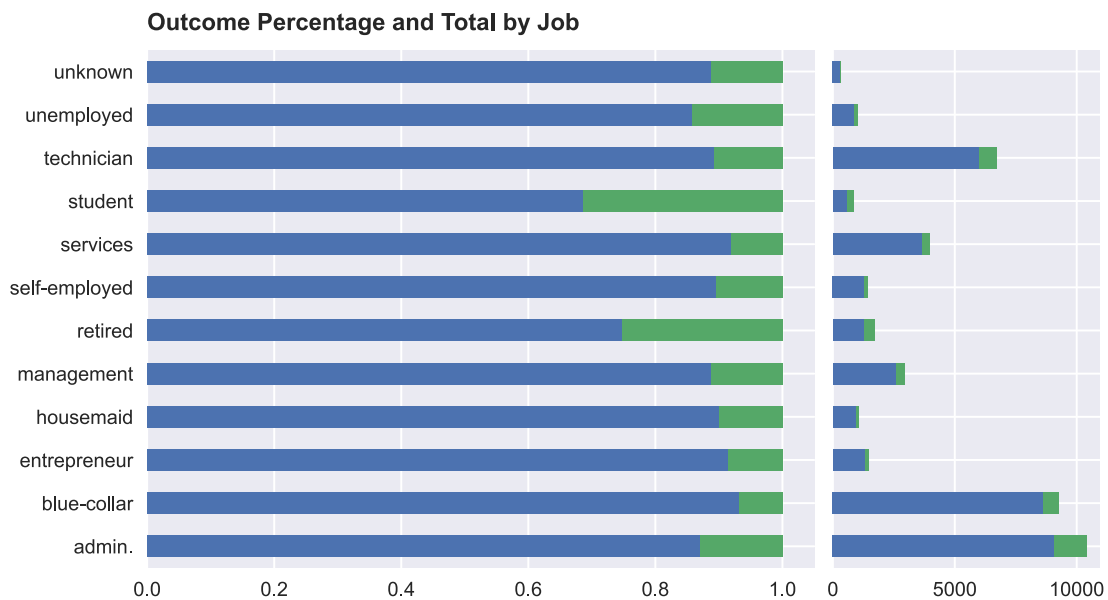
```
[21]: def cat_outcome(df, feature):
    df = df.copy()
```

```

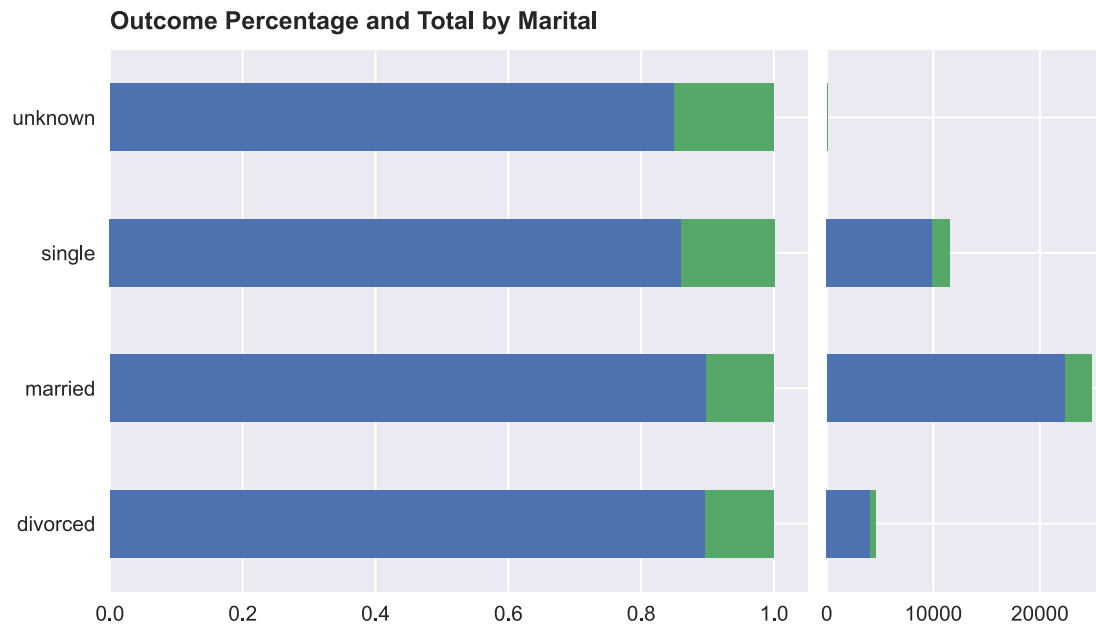
    if pd.api.types.is_categorical_dtype(df[feature]) and df[feature].isna().
    sum() > 0:
        df[feature] = df[feature].cat.add_categories("unknown")
        df[feature] = df[feature].fillna("unknown")
        title = feature.title().replace("_", " ").replace("Of", "of")
        f, axs = plt.subplots(1, 2, figsize=(8.6, 4.8), sharey=True,
    gridspec_kw=dict(wspace=0.04, width_ratios=[5, 2]))
        ax0 = df["y"].groupby(df[feature], dropna=False).
    value_counts(normalize=True).unstack().plot.barh(xlabel="", legend=False,
    stacked=True, ax=axs[0], title=f"Outcome Percentage and Total by {title}")
        ax1 = df["y"].groupby(df[feature], dropna=False).value_counts().unstack().
    plot.barh(xlabel="", legend=False, stacked=True, ax=axs[1])

```

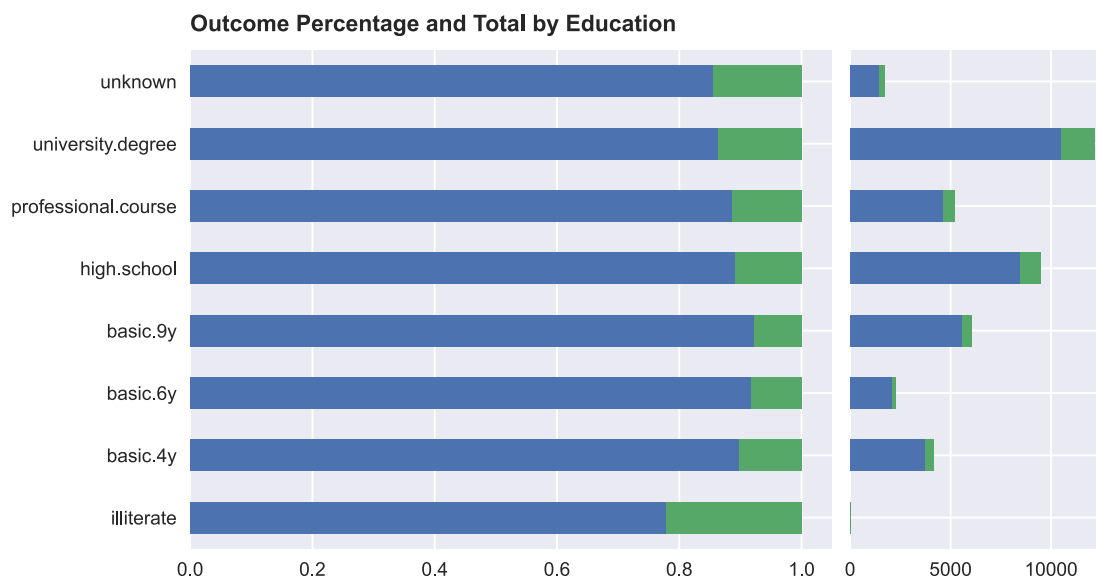
```
[22]: job_outcome = cat_outcome(bank_mkt, "job")
```



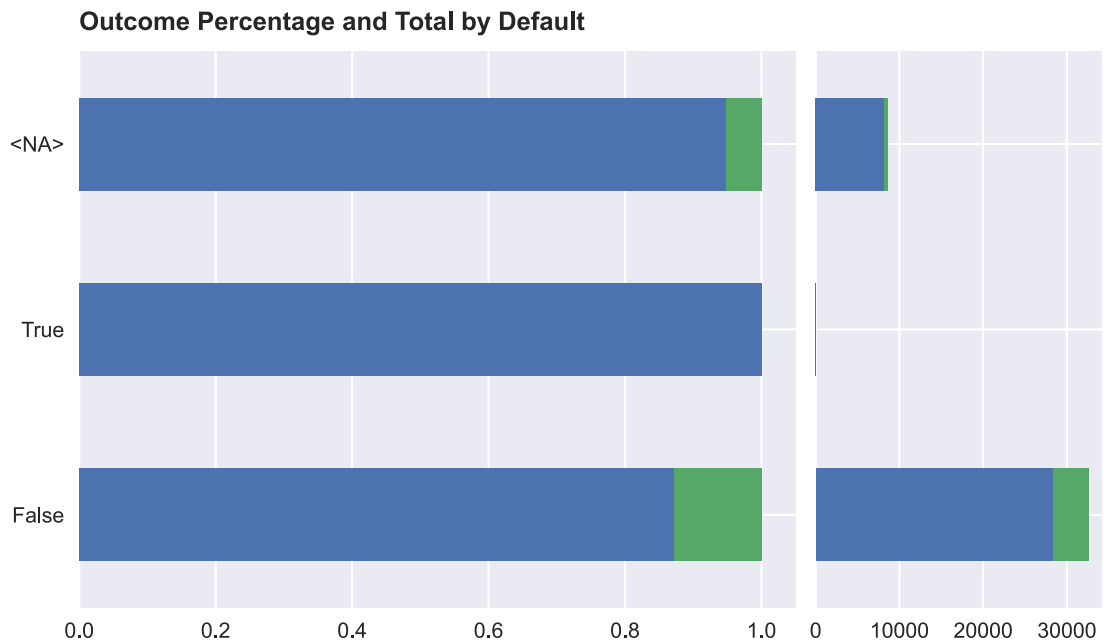
```
[23]: marital_outcome = cat_outcome(bank_mkt, "marital")
```


```
[24]: education_outcome = cat_outcome(bank_mkt, "education")
```




```
[25]: default_outcome = cat_outcome(bank_mkt, "default")
```




```
[26]: housing_outcome = cat_outcome(bank_mkt, "housing")
```




```
[27]: loan_outcome = cat_outcome(bank_mkt, "loan")
```




```
[28]: job_marital_total = bank_mkt[["job", "marital", "y"]].groupby(["job",
    ↳ "marital"]).count().y.unstack()
job_marital_true = bank_mkt[["job", "marital", "y"]].groupby(["job",
    ↳ "marital"]).sum().y.unstack()
job_marital_rate = job_marital_true / job_marital_total
job_marital_rate = job_marital_rate.rename_axis(None, axis=0).rename_axis(None,
    ↳ axis=1)
job_marital_heatmap = sns.heatmap(data=job_marital_rate, vmin=0, vmax=0.5,
    ↳ annot=True).set_title("True Outcome Percentage by Job and Marital Status")
```



```
[29]: job_education_total = bank_mkt[["job", "education", "y"]].groupby(["job", "education"]).count().y.unstack()
      ↪ "education"]
job_education_true = bank_mkt[["job", "education", "y"]].groupby(["job", "education"]).sum().y.unstack()
      ↪ "education"]
job_education_rate = job_education_true / job_education_total
job_education_rate = job_education_rate.rename_axis(None, axis=0).
      ↪ rename_axis(None, axis=1)
job_education_heatmap = sns.heatmap(data=job_education_rate, vmin=0, vmax=0.5,
      ↪ annot=True).set_title("True Outcome Percentage by Job and Education")
```




```
[30]: education_marital_total = bank_mkt[["education", "marital", "y"]].
      ↪groupby(["education", "marital"]).count().y.unstack()
education_marital_true = bank_mkt[["education", "marital", "y"]].
      ↪groupby(["education", "marital"]).sum().y.unstack()
education_marital_rate = education_marital_true / education_marital_total
education_marital_rate = education_marital_rate.rename_axis(None, axis=0).
      ↪rename_axis(None, axis=1)
education_marital_heatmap = sns.heatmap(data=education_marital_rate, vmin=0,
      ↪vmax=0.5, annot=True).set_title("True Outcome Percentage by Education and
      ↪Marital Status")
```

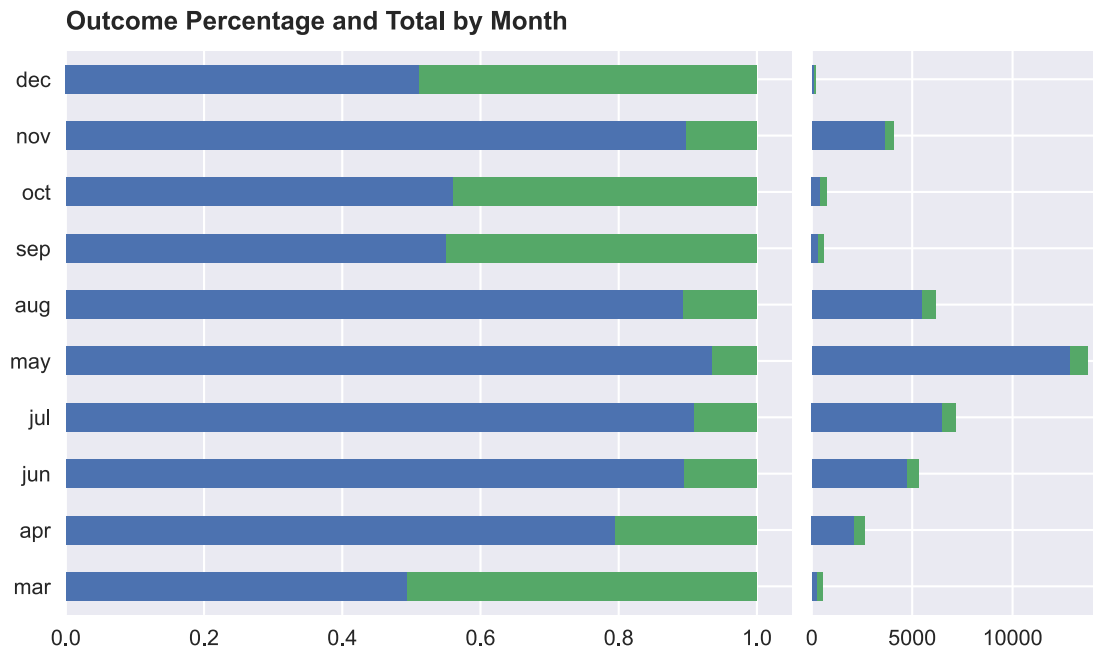


1.2.3 Current Campaign


```
[31]: contact_outcome = cat_outcome(bank_mkt, "contact")
```



```
[32]: month_outcome = cat_outcome(bank_mkt, "month")
```



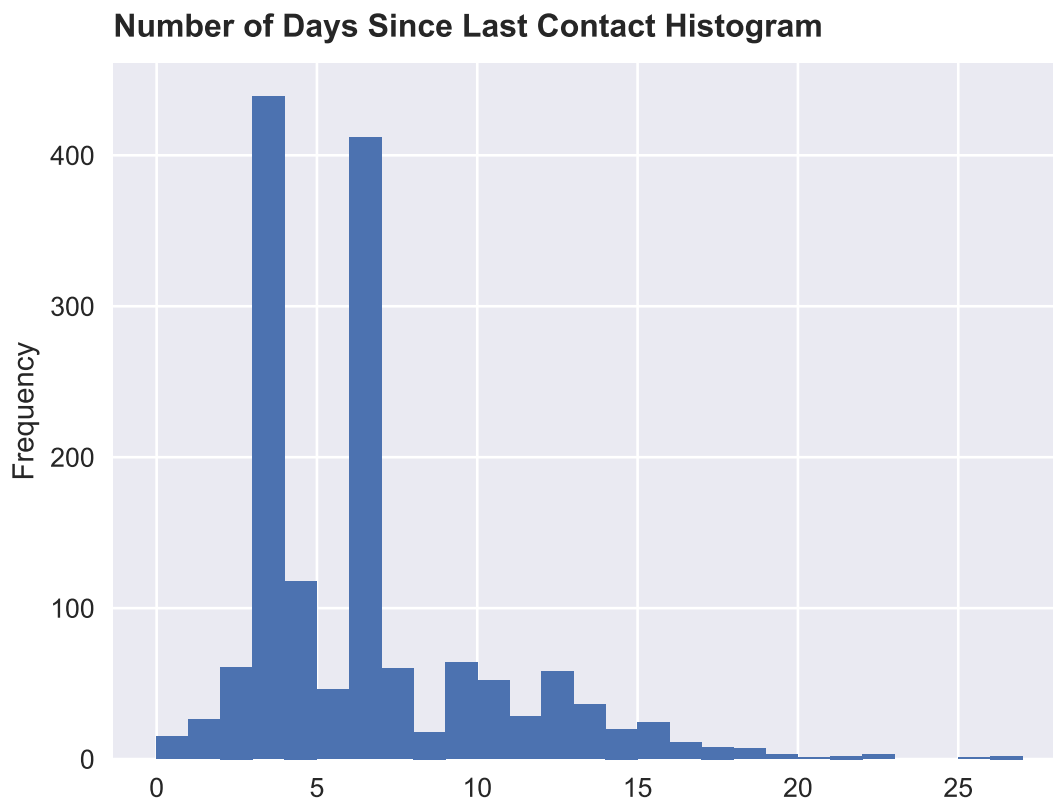
```
[33]: day_outcome = cat_outcome(bank_mkt, "day_of_week")
```



1.2.4 Previous Campaign

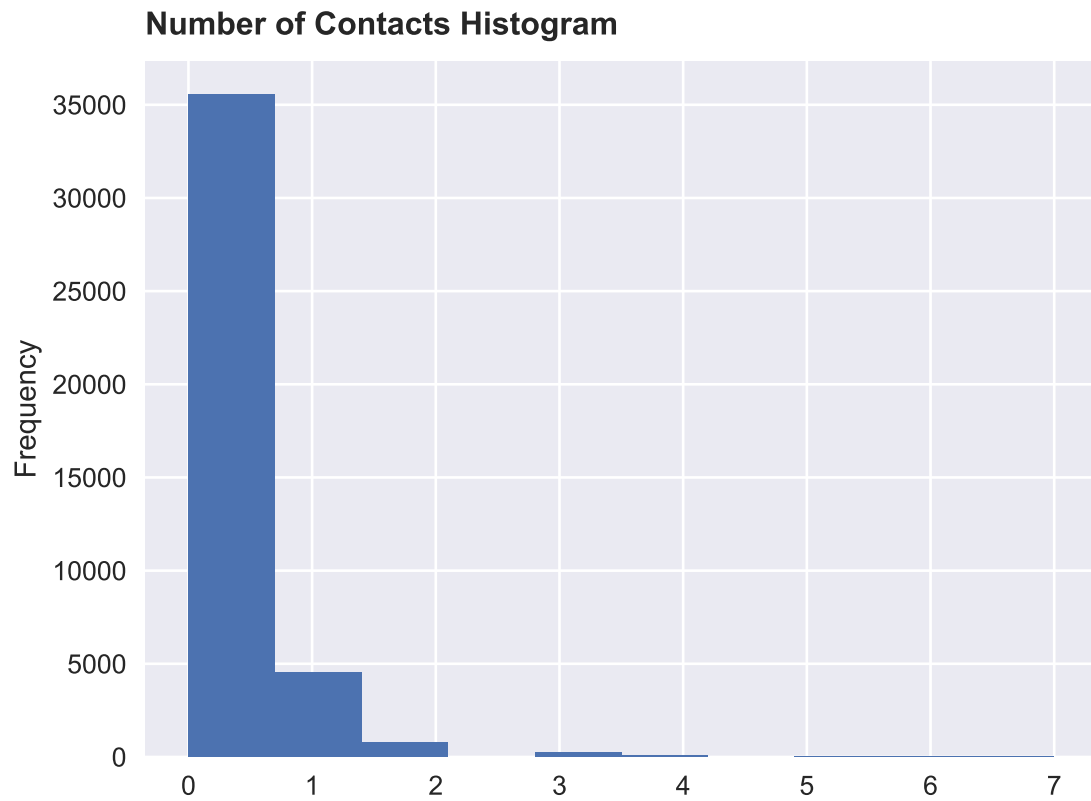
We can plot the distribution of `pdays` and `previous`. As we can see, most of the client with `pdays` has been contacted 3 to 6 days before and peaked at 3 and 6 days.

```
[34]: pdays_hist = bank_mkt["pdays"].plot.hist(bins=27, title="Number of Days Since_␣  
      ↪Last Contact Histogram")
```




Most of the client has never been contacted before.

```
[35]: previous_hist = bank_mkt["previous"].plot.hist(title="Number of Contacts_␣  
      ↪Histogram")
```

If `pdays` is missing value, that means that the client was not previously contacted and therefore should not have `poutcome`. But `poutcome` column has less missing values than `pdays`.

```
[36]: previous_na = bank_mkt[["pdays", "poutcome"]].isna().sum()
      previous_na_ax = previous_na.plot.bar(title="Number of Missing Values in pdays_
      ↪and poutcome")
```



We can print out the 4110 rows where the client is not contacted but have **poutcome** and see how many times they have been contacted before. The figures suggest that maybe these clients has been actually contacted but it was more than 30 days ago so the contact date was not recorded. This leaves us plenty room for feature engineering.

```
[37]: previous = bank_mkt[["campaign", "pdays", "previous", "poutcome", "y"]]
previous = previous[previous["pdays"].isna() & previous["poutcome"].notna()]
↪ reset_index(drop=True)
previous
```


```
[37]:
```

	campaign	pdays	previous	poutcome	y
0	1	<NA>	1	False	False
1	1	<NA>	1	False	True
2	1	<NA>	1	False	False
3	1	<NA>	1	False	True
4	1	<NA>	1	False	False
...
4105	1	<NA>	1	False	True
4106	2	<NA>	4	False	False

4107	1	<NA>	2	False	True
4108	1	<NA>	2	False	False
4109	3	<NA>	1	False	False

[4110 rows x 5 columns]

```
[38]: previous_ax = previous["previous"].plot.hist(bins=12, title="Number of Days Since Last Contact for Previous Client Histogram")
```



```
[39]: bank_mkt[bank_mkt["pdays"].isna() & bank_mkt["poutcome"].isna()]
```

```
[39]:
```

	age	job	marital	education	default	housing	\
0	56	housemaid	married	basic.4y	False	False	
1	57	services	married	high.school	<NA>	False	
2	37	services	married	high.school	False	True	
3	40	admin.	married	basic.6y	False	False	
4	56	services	married	high.school	False	False	
...	
41181	37	admin.	married	university.degree	False	True	
41183	73	retired	married	professional.course	False	True	
41184	46	blue-collar	married	professional.course	False	False	
41185	56	retired	married	university.degree	False	True	

41186	44	technician	married	professional.course	False	False
-------	----	------------	---------	---------------------	-------	-------

	loan	contact	month	day_of_week	...	campaign	pdays	previous	\
0	False	telephone	may	mon	...	1	<NA>	0	
1	False	telephone	may	mon	...	1	<NA>	0	
2	False	telephone	may	mon	...	1	<NA>	0	
3	False	telephone	may	mon	...	1	<NA>	0	
4	True	telephone	may	mon	...	1	<NA>	0	
...	
41181	False	cellular	nov	fri	...	1	<NA>	0	
41183	False	cellular	nov	fri	...	1	<NA>	0	
41184	False	cellular	nov	fri	...	1	<NA>	0	
41185	False	cellular	nov	fri	...	2	<NA>	0	
41186	False	cellular	nov	fri	...	1	<NA>	0	

	poutcome	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	\
0	<NA>	1.1	93.994	-36.4	4.857	
1	<NA>	1.1	93.994	-36.4	4.857	
2	<NA>	1.1	93.994	-36.4	4.857	
3	<NA>	1.1	93.994	-36.4	4.857	
4	<NA>	1.1	93.994	-36.4	4.857	
...	
41181	<NA>	-1.1	94.767	-50.8	1.028	
41183	<NA>	-1.1	94.767	-50.8	1.028	
41184	<NA>	-1.1	94.767	-50.8	1.028	
41185	<NA>	-1.1	94.767	-50.8	1.028	
41186	<NA>	-1.1	94.767	-50.8	1.028	

	nr.employed	y
0	5191.0	False
1	5191.0	False
2	5191.0	False
3	5191.0	False
4	5191.0	False
...
41181	4963.6	True
41183	4963.6	True
41184	4963.6	False
41185	4963.6	False
41186	4963.6	True

[35563 rows x 21 columns]

1.2.5 Correlation Heatmap

```
[40]: corr_heatmap = sns.heatmap(data=bank_mkt.corr(method="pearson")).  
      ↪set_title("Correlation Heatmap")
```

